PCET-NMVPM's

Nutan College of Engineering and Research, Talegaon, Pune

PROJECT FILE

[AY: 2025-2026]

Project Domain: Artificial Intelligence

Project Title: Jurisynth: AI for Legal Clarity Efficiency

Project Group Member:

Prathamesh Sachin Holay	**[23064191242506]**
Gargi Prashant Thipase	**[23064191242509]**
Shrikant Chaturbhuj Jadhav	**[23064191242511]**
Pranita Santosh Panchal	**[23064191242515]**

Project Guide: Prof. Jordan Choudhari

Department of Computer Science and Engineering

A

PROJECT PHASE I

REPORT ON

"Jurisynth: AI for Legal Clarity Efficiency"

Submitted to the

Dr. Babasaheb Ambedkar Technological University Lonere, Raigad

in fulfillment of the requirements

for the award of the degree

BACHELORS OF TECHNOLOGY

COMPUTER SCIENCE AND ENGINEERING
2025-2026

BY

Prathamesh Sachin Holay	**[23064191242506]**
Gargi Prashant Thipase	**[23064191242509]**
Shrikant Chaturbhuj Jadhav	**[23064191242511]**
Pranita Santosh Panchal	**[23064191242515]**

UNDER THE GUIDANCE OF

Prof. Jordan Choudhari

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

PCET-NMVPM's

NUTAN COLLEGE OF ENGINEERING AND RESEARCH

TALEGAON, PUNE 410507

PCET-NMVPM's

NUTAN COLLEGE OF ENGINEERING & RESEARCH TALEGAON, PUNE

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

CERTIFICATE

This is to certify that the Project Report entitled **"Jurisynth: AI for Legal Clarity Efficiency",** which is being submitted by**, Prathamesh Holay, Gargi Thipase, Shrikant Jadhav, Pranita Panchal** as partial fulfillment for the **Degree Bachelor of Technology** (Computer Science and Engineering) of **DBATU, Lonere.**

This is bonafide work carried under my supervision and guidance.

Place: Pune

Date:

Prof. Jordan Choudhari
Project Guide

Prof. Shital Mehta
Project I/C

Dr. Mahesh Wankhade
Head of Department

Dr. Pramod Patil
Principal

External Examiner [Name & Sign]

SEAL

ACKNOWLEDGEMENT

The satisfaction and euphoria that accompany the successful completion of any task would be incomplete without the mention of people who made it possible. So, we acknowledge all those whose guidance and encouragement served as a beacon light and crowned our efforts with success.

We have immense pleasure in expressing thanks to the principal ***Dr. Pramod Patil*** for providing all the facilities for the successful completion of the project.

With due respect, we thank our Head of Department ***Dr. Mahesh Wankhade, Computer Science and Engineering***, for his motivating support, keen interest which kept our spirits alive all through.

We would like to express thanks to our guide ***Prof. Jordan Choudhari*** Department of ***Computer Science and Engineering*** who has guided us throughout the completion of this project.

Finally, we would like to thank ***all the teaching and non-teaching staff and all our friends*** who have rendered their support in the completion of this report.

Prathamesh Sachin Holay [23064191242506]
Gargi Prashant Thipase [23064191242509]
Shrikant Chaturbhuj Jadhav [23064191242511]
Pranita Santosh Panchal [23064191242515]

ABSTRACT

The report titled **"Jurisynth: AI for Legal Clarity and Efficiency"** examines how **Artificial Intelligence (AI)** and **Natural Language Processing (NLP)** can reshape the legal ecosystem by simplifying, summarizing, and interpreting complex legal documents. Traditionally, the legal field has relied on expert human judgment to analyze, interpret, and summarize case laws, contracts, and statutory materials. However, the explosive increase in digital legal data has rendered manual review slow, inefficient, and prone to inaccuracies. **Jurisynth** addresses this challenge by introducing an automated framework designed to simplify and summarize legal texts while maintaining their semantic integrity.

The proposed system combines **lexical simplification**, **syntactic restructuring**, and **hybrid summarization** using state-of-the-art transformer models such as **Legal-BERT** and **Legal-Pegasus**. A **human-in-the-loop** validation mechanism ensures that legal experts can evaluate AI-generated summaries for accuracy, reliability, and ethical compliance.

This report further discusses the system architecture, methodology, and projected outcomes, supported by an extensive literature survey on modern legal NLP techniques. Key focus areas include **ethical AI design**, fairness, accountability, and transparency—crucial factors in high-stakes legal applications.

The ultimate vision of **Jurisynth** is to promote AI-assisted access to justice by transforming complex legal language into clear, precise, and contextually consistent summaries. By bridging the gap between intricate jurisprudence and everyday citizens, **Jurisynth** advances both legal inclusivity and technological innovation.

LIST OF CONTENTS

INDEX

LIST OF FIGURES

LIST OF TABLES

LIST OF ABBREVIATIONS

Sr. No	Abbreviations	Description
1	**AI**	Artificial Intelligence
2	**NLP**	Natural Language Processing
3	**ML**	Machine Learning
4	**LLM**	Large Language Model
5	**BERT**	Bidirectional Encoder Representations from Transformers

CHAPTER 1 : INTRODUCTION

1.1 About

In today's fast-evolving legal ecosystem, the demand for accuracy, speed, and clarity has never been greater. Legal professionals, researchers, and clients alike face increasing challenges in accessing relevant laws, interpreting complex legal texts, and managing large volumes of case data. Traditional methods of legal research are often time-consuming, resource-intensive, and prone to human oversight. To address these challenges, Jurisynth emerges as an innovative AI-powered solution designed to streamline legal research, enhance comprehension, and elevate overall legal efficiency.

Jurisynth is built to bridge the gap between complex legal information and user understanding. Leveraging advanced natural language processing (NLP) and machine learning techniques, the system enables users to search, retrieve, and interpret legal content with greater precision. It offers a unified interface through which lawyers, law students, researchers, and even general users can access case references, statutory sections, and legal explanations quickly and accurately. By automating tedious research tasks and providing structured, section-wise insights, Jurisynth significantly reduces manual workload while improving the quality of legal interpretation. A key strength of Jurisynth lies in its hybrid architecture. It integrates API-based legal data retrieval for up-to-date legal sections alongside custom AI models capable of generating simplified explanations, summaries, and context-aware responses. This combination ensures both reliability and flexibility: users receive authenticated legal content complemented by dynamic AI-generated clarity. The platform further supports multi-format inputs, including text and audio, allowing users to upload voice queries and receive translated or transcribed outputs—particularly valuable for users unfamiliar with legal jargon or technical language. To ensure secure access, Jurisynth implements a multi-layer authentication system, including system-generated IDs, CAPTCHA verification, and optional Google Login for convenience. This helps maintain the confidentiality and integrity of legal information, especially when used by law firms or professional groups. Ultimately, Jurisynth aims to redefine how legal information is accessed and understood. By combining the precision of structured legal databases with the interpretability and adaptability of AI models, it enhances clarity, reduces research time, and empowers users to make informed decision

1.2 Necessity

The necessity for Jurisynth arises from the growing demand for clarity, speed, and accuracy within the legal ecosystem. As legal systems expand and evolve, the volume of statutes, case laws, amendments, and judicial interpretations has increased exponentially. This creates a significant challenge for lawyers, law students, researchers, and general users who must navigate this vast information landscape. Traditional legal research methods rely heavily on manual reading, cross-referencing, and interpretation, making the process slow, labor-intensive, and vulnerable to human oversight. Jurisynth becomes essential in addressing these limitations by offering an AI-driven mechanism that drastically improves the efficiency and reliability of legal research.

The legal domain is inherently complex, relying on vast bodies of statutes, case laws, regulations, and interpretations that continually evolve. As a result, legal professionals, students, and individuals often struggle to access accurate information quickly and interpret it effectively. The necessity for a system like Jurisynth arises from multiple challenges that limit productivity, accuracy, and accessibility within the current legal research and interpretation process.

Firstly, traditional legal research is time-consuming and labor-intensive. Lawyers and researchers are required to manually read through extensive case files, judgments, and legislative documents to find relevant information. This slows down the decision-making process and increases the risk of overlooking crucial details. Jurisynth streamlines this workflow by automating search, classification, and summarization, significantly reducing research time.

Secondly, legal language is often dense, technical, and difficult to understand for non-experts. Clients and general users frequently face difficulties interpreting legal terms, procedures, and sections. Jurisynth addresses this by providing simplified, AI-generated explanations and summaries that enhance comprehension without compromising accuracy. Another critical necessity stems from the lack of centralized and organized access to legal content. Legal documents may be scattered across multiple databases, formats, and sources. Jurisynth consolidates this fragmented ecosystem by integrating API-driven legal data and AI-powered interpretation into a single platform, ensuring faster and easier accessibility.

Moreover, the legal sector is witnessing increasing demand for digital transformation and automation. Law firms and legal departments aim to improve operational efficiency, minimize repetitive tasks, and enhance service quality. Jurisynth supports this shift by automating routine research functions, enabling professionals to dedicate more time to strategy, analysis, and client consultation. Accessibility also remains a significant issue, especially for individuals in rural or underserved areas. Many users struggle with language barriers or limited access to professional legal assistance. Jurisynth's support for audio queries, translations, and simplified outputs ensures that legal information becomes more inclusive and user-friendly.

Lastly, maintaining accuracy and ensuring up-to-date information is crucial. Manual legal research runs the risk of relying on outdated interpretations or missing recent amendments. By integrating automated data retrieval, Jurisynth ensures that users access the most current legal sections, judgments, and interpretations available.

1.3 Problem Statement

II. Contemporary legal information platforms remain limited in functionality as they primarily provide predefined content without offering contextual interpretation or multimodal data processing. The inability to analyze audio inputs, extract structured data from uploaded documents, or generate section-specific insights restricts their usefulness. A comprehensive AI-driven system is required to enhance legal comprehension through real-time summarization, automated interpretation, and secure access.

III. Legal work also requires quick interpretation of sections, extraction of text from documents, and generation of clear summaries, which becomes easier when AI can process audio, PDFs, images, and online references together. By delivering real-time clarity and organized insights, such a system can support faster understanding and more efficient decision-making for users.

1.4 Motivation

The motivation behind developing an AI-driven legal platform comes from the growing need to make legal information more understandable, accessible, and faster to work with. Lawyers, students, and everyday users often deal with multiple sources books, online references, case documents, and audio conversations making it difficult to gather and interpret information efficiently. By creating a system that can extract text from documents, process audio inputs, understand sections, and summarize content in a clear manner, we aim to simplify the entire legal research experience.

The project is driven by a desire to combine accuracy with convenience, allowing users to receive instant clarity without searching through lengthy documents or interpreting complex legal terms.

Integrating hybrid AI models with legal APIs enables the platform to deliver context-aware insights, making legal understanding smoother and more intuitive. This approach not only saves time but also supports better decision-making by presenting organized, easy-to-follow legal information tailored to the user's query.

1.5 Objective

1. To review existing AI-driven legal information systems and multimodal document-processing techniques for identifying gaps in automated legal analysis.

2. To design and develop an integrated AI framework capable of extracting, structuring, and interpreting legal information from text, documents, images, and audio queries.

3. To evaluate the accuracy, clarity, and usability of the proposed Jurisynth system in enhancing legal understanding and decision-making efficiency.

CHAPTER 2: LITERATURE SURVEY

2.1 Literature Survey

Sr. No.	Author / Year	Title of Paper	Technology Used	Description	Limitations
1	Nigam et al., 2024	NyayaAnumana & INLegalLlama	Legal LLMs, Indian Judgment Prediction	Dataset + model designed for Indian judgment prediction with improved accuracy on reasoning tasks.	Focused mainly on higher courts; requires heavy compute.
2	Nigam et al., 2025	TathyaNyaya & FactLegalLlama	Factual LLMs, Explainable AI	Factual judgment prediction with structured explanations based on case facts.	Heavily dependent on quality of fact extraction.
3	Nigam et al., 2024	PredEx	Explainable ML, Judgment Interpretation	Provides interpretable predictions aligned with judicial reasoning.	Explanations may not fully reflect actual court logic.

4	Nigam et al., 2025	NyayaRAG	Retrieval-Augmented Generation (RAG)	Combines legal retrieval with AI reasoning for grounded predictions.	Retrieval accuracy depends on embedding quality.
5	Parikh et al., 2021	LawSum	Weakly Supervised Summarization	Generates summaries using noisy supervision from headnotes.	Weak labels cause inconsistency and occasional missing context.
6	Datta et al., 2023	MILDSum	Multilingual Summarization	Enables summarization across multiple Indian languages.	Performance varies on low-resource regional languages.
7	Shukla et al., 2022	Legal Case Document Summarization	Extractive + Abstractive NLP	Examines hybrid approaches for judgment summarization.	Abstractive models risk hallucinations.

8	Kalamkar et al., 2022	Legal Document Structuring	Rhetorical Role Labeling	Segments text into facts, arguments, reasoning, decision.	Fails on non-standard court formatting.
9	Pallavi et al., 2025	LEGAL AI	Search + Prediction Engine	Basic legal research + prediction system for courts.	Limited depth; not optimized for complex cases.
10	Verma & Singh, 2023	Explainable Legal AI	Explainability Frameworks	Provides interpretable judgment prediction models.	Less accurate than large-scale LLMs.
11	Malik et al., 2021	ILDC for CJPE	Indian Legal Dataset	Large labeled dataset for case prediction and explanation.	Labels contain some noise; domain imbalance exists.

12	Trivedi et al., 2024	ILC (Indian Legal Corpus)	Text Cleaning + NLP	Large corpus of Indian legal proceedings.	Still contains court formatting inconsistencies.
13	Sivaranjani et al., 2023	Supreme Court Dataset	Outcome Classification	Predictive dataset for Supreme Court outcomes.	Limited generalization to lower courts.
14	Goswami et al., 2024	Legal Summaries with Domain Knowledge	Multi-Objective Optimization	Improves legal summary informativeness via domain cues.	High complexity; requires tuning.
15	Nigam et al., 2025	LegalSeg	Text Segmentation	Segments judgments into rhetorical roles.	Accuracy drops for poor-quality PDFs.

16	Paul et al., 2022	InLegalBERT	Domain Pretrained BERT	Legal-domain BERT for Indian judgments.	Short context window limits long-case handling.
17	Sharma & Singh, 2024	InLegalBERT Summarizer	Transformer Summarization	Optimized for long Indian judgments summarization.	Struggles with extremely long documents.
18	Prabhakar & Pati, 2024	Extractive Legal Summaries	Hybrid Extractive-Abstractive	Safe summaries using extraction + generative refinement.	Limited creativity; depends on extractive boundaries.
19	Kumar, 2025	Summarization via Section Classification	Classification + Extraction	Focuses summary on judgment-critical sections.	Requires well-labeled training data.

20	Pati et al., 2025	LegalSummNet	Long-Context Transformer	Advanced summarizer for long legal documents.	High GPU and memory requirements.
21	Survey, 2025	Indian Legal Summaries Survey	Comparative Study	Compares summarization techniques for legal text.	Contains no new dataset or model.
22	Advancements in Legal NLP, 2023	Law NLP Review	Survey + Analysis	Overview of ML/NLP for legal tasks.	Does not focus deeply on Indian domain issues.
23	Zahra et al., 2025	Legal Document Summarizer	Template + Neural Hybrid	Combines template safety with neural fluency.	Template constraints limit flexibility.

24	IndicLegalQA, 2024	Legal QA Dataset	Retrieval + QA Models	Provides QA pairs over Indian judgments.	Limited dataset size.
25	Mahapatra et al., 2025	MILPaC	Legal Machine Translation	Benchmark for translating Indian legal acts.	Domain terminology errors are common.
26	Aumiller et al., 2022	EUR-Lex-Sum	Cross-Lingual Summarization, Long-Form NLP	Large EU legal summarization dataset with multilingual long-text processing.	Not tailored for Indian legal system.
27	Xiao et al., 2018	CAIL2018	Judgment Prediction, Large Legal Dataset	One of the largest Chinese legal datasets for prediction and classification.	Different legal system; limited transfer to common-law structure.

28	Huang et al., 2023	Two-Stage Legal Summarization	Two-Stage Transformer Summaries	High-precision summarization via extraction → abstraction pipeline.	Struggles with highly complex long judgments.
29	Krishna & Reddy, 2019	Legal Summarization using GSA	Swarm Intelligence, Extractive NLP	Optimization-based extraction using Gravitational Search Algorithm.	Extraction-only; no abstractive capability.
30	Shang, 2022	Computational Intelligence for Legal Prediction	Computational Intelligence Models	Framework using ML for decision support in courts.	Lacks long-context encoding; shallow ML models.
31	Richmond et al., 2023	Explainable AI & Law	Explainable AI Survey	Evidential analysis of XAI techniques in law.	Survey only; no implementation or dataset.

32	Mentzingen et al., 2025	Cost-Efficient Legal Precedent Retrieval	LLMs + Summarization + Retrieval	Aims to reduce inference cost while improving precedent recall.	Still not optimized for very long judgments.
33	Gersh et al., 2021	CUAD Dataset	Contract Understanding, NLP	Large contract review dataset for clause extraction.	Focuses on contracts, not judgments.
34	Pallavi et al., 2025	LEGAL AI (IJARCCE)	AI Search + Case Prediction	Legal research system integrating search + prediction.	Basic design; lacks deep reasoning pipeline.
35	Verma & Singh, 2023	Explainable Legal Prediction	Explainable ML, SHAP/Attention	Transparent decision-making for court judgment prediction.	Lower accuracy than transformer models.

36	Kaczmarek et al., 2021	Atticus Contract Understanding	Contract NLP, Clause Extraction	High-quality contract clause extraction dataset.	Non-Indian domain; contract-only focus.
37	Heddaya et al., 2025	CaseSumm Dataset	Long-Context Summarization	U.S. Supreme Court long-doc summarization dataset.	Not Indian; different court structure and language.
38	Pati & Ghosh, 2024	Indian Legal RAG Systems	RAG, Vector Retrieval, LLM	Studies design of retrieval-augmented systems for Indian law.	Performance depends on DB density and embedding quality.
39	Thomas et al., 2024	Legal Judgment Explanation Models	Explanation Generation, NLP	Generates interpretive explanations for model outputs.	May produce generic explanations.

40	Sharma et al., 2025	Legal Fact Extraction	Model-Based Fact Extraction	Extracts fact segments from judgments before prediction.	Struggles on noisy input PDFs.
41	Ghosh et al., 2023	Case Law Retrieval Benchmarks	Retrieval Models, Dense Passage Retrieval	Benchmark for retrieving similar Indian legal cases.	Does not include modern LLM-based evaluation.
42	Arora et al., 2022	Indian Court Document OCR	OCR + NLP Pipeline	Improves text extraction quality from scanned judgments.	Accuracy depends on scan quality.
43	Pradhan et al., 2024	Legal Argument Mining	Argument Mining, NLP	Extracts reasoning, arguments, and issue statements.	Argument segmentation remains imperfect.

44	Mishra et al., 2023	AI for Bail Decision Support	ML Classification	Supports bail decision-making using structured ML.	Controversial due to bias concerns.
45	Shallum et al., 2024	Legal Long-Context Models	Long-Context Transformers	Optimizes LLMs for extremely long legal documents.	Very high GPU memory requirements.
46	Patnaik et al., 2025	Judicial Outcome Modeling	Statistical + Neural Prediction	Combines statistical signals with deep models for robust predictions.	Requires large labeled datasets.
47	Reddy et al., 2023	Legal Topic Classification	Topic Modeling + BERT	Classifies legal text into multi-label legal topics.	Struggles with overlapping classes.

48	Chouhan et al., 2024	Cross-Jurisdiction Legal NLP	Cross-Lingual Models	Trains multilingual legal models across jurisdictions.	Differences in legal systems limit performance.
49	Gupta et al., 2025	Legal Document Embeddings	Law-Specific Embedding Models	Builds embeddings optimized for legal similarity search.	Embedding drift over long documents.
50	Khan et al., 2025	Judgment Summaries via LLMs	Transformer LLMs	Summarizes long judgments with improved coherence and readability. occasional hallucination in generative output.	

CHAPTER 2.1: GAP IDENTIFICATION

❖ Limitations of Current AI Research in the Indian Legal Domain

1. Fragmented and Single-Task Systems

- Most AI models focus on only one task (judgment prediction, summarization, or document retrieval).
- They fail to support complete legal workflows such as legal reasoning, drafting assistance, statutory mapping, or decision-support.
- No holistic platform integrates interpretation, summarization, advisory recommendations, and precedent justification.

2. Inadequate and Unrepresentative Datasets

- Public datasets mostly include Supreme Court or foreign judgments.
- Trial court cases, regional court data, and multilingual content are poorly represented.
- Models suffer from bias, data imbalance, lack of linguistic diversity, and failure to capture India's legal complexity.

3. Lack of Explainability and Transparency

- Many predictive models provide accuracy but not legal reasoning or logical justification.
- Absence of transparent explanations reduces trust among judges, lawyers, and legal practitioners.
- Evaluations rely heavily on metrics like ROUGE or BLEU, with limited expert involvement.

4. Poor Practical Usability

- Most proposed systems remain academic prototypes without deployment-ready architecture.
- Missing components include privacy safeguards, security frameworks, user-friendly interfaces, and ethical controls.
- Issues like hallucination, incorrect suggestions, and lack of accountability hinder real-world use.

5. Limited Adoption and Domain Acceptance

- Current tools do not align with the practical needs of advocates, judicial officers, or citizens seeking legal support.
- Lack of user-oriented design creates a gap between AI capabilities and real courtroom workflows.

CHAPTER 3 : PROPOSED SYSTEM

3.1 Tech Stack

The proposed system, **JuriSynth**, is built as an AI-powered legal analysis platform designed to summarize judgments, retrieve precedents, and generate judicial predictions using long-context transformer models and retrieval-augmented generation (RAG).

The tech stack integrates a modern web-based UI, secure backend APIs, domain-tuned legal NLP models, and vector-based retrieval for scalable legal intelligence.

It uses a combination of **Node.js / Python**, **LLM model servers**, and **vector databases (FAISS / Pinecone/ChromaDB)** to efficiently process long Indian legal documents.
The system is deployed using a microservices-style architecture for reliability, performance, and modularity.

3.1.1 The languages used :

Python
Used for implementing the NLP pipeline, including:

- Judgment summarization
- Precedent retrieval
- Feature extraction
- Tokenization & embedding
- Running transformer/LLM models
- Preprocessing PDFs and long-text formatting

Python provides strong support through libraries such as **Transformers, LangChain, HuggingFace, spaCy, PyMuPDF**, making it ideal for training and deploying legal AI models.

JavaScript/TypeScript
Used in the backend (Node.js) and frontend (React):

- Backend APIs for serving model results
- Authentication & role-based access
- User dashboards for lawyers
- Communicating with AI inference endpoints
- Displaying summaries, predictions, and legal citations

GQL/JSON/RESTAPIs
For system-to-system communication between frontend, backend, and model servers.

3.1.2 Software Requirement

Backend: Node.js / Express Server

Handles:

- User authentication (lawyer, student, admin)
- Routing requests to AI model server
- Handling PDFs and text extraction
- Logging queries and storing summaries
- Secure session-based interactions

Provides middleware for validating input, managing request flow, and maintaining access logs.

AI Processing Environment (Python)

Includes:

- Preprocessing pipeline
- Transformer-based summarization model
- Indian legal judgment prediction model
- RAG pipeline for retrieving relevant precedents
- Vector database interface

Libraries used:

- **HuggingFace Transformers**
- **LangChain**
- **FAISS / ChromaDB**
- **PyMuPDF / PDFminer**
- **SentencePiece tokenization**

Database Layer

The system uses:

- **PostgreSQL / MongoDB** for storing users, queries, summaries, predictions
- **Vector Database (FAISS / Pinecone)** for legal document embeddings
- **S3 / Local storage** for uploaded PDFs

Frontend (React / Next.js)

Used for:

- Lawyer dashboard
- Uploading judgments
- Displaying AI-generated summaries
- Viewing relevant case laws
- Predictions with explanations
- Downloadable reports

3.1.3 Hardware Requirement

This project is software-centric, but it requires computational infrastructure for running large language models efficiently.

Model Inference Hardware

- **GPU Server (NVIDIA 12–40 GB VRAM)** for running long-context models
- OR **Cloud inference API** (OpenAI, HuggingFace Inference, custom GPU node)

Local Machine Requirements (for development)

- Minimum 8 GB RAM (16 GB recommended)
- Python 3.10+
- Node.js 18+
- Storage for datasets & embeddings (5–15 GB)

Server Requirements

- Ubuntu 22.04 LTS
- Docker / Docker Compose
- HTTPS + certificate
- Reverse proxy (NGINX)
- Firewall, security hardening

These requirements ensure scalable, secure, and real-time AI inference for legal analysis.

3.1.4 System Overview

JuriSynth is an **AI-powered legal research and judgment-support system** capable of analyzing long legal texts and producing:

- Structured judgment summaries
- Extracted facts, issues, arguments
- Relevant precedent suggestions
- Legal issue segmentation
- Judgment outcome prediction
- Explanation for predictions

The system integrates **OCR, NLP, transformer models, embeddings, retrieval, and reasoning modules** into a single pipeline.

System Workflow

1. **User uploads a judgment or enters text.**
2. The backend extracts text, cleans formatting, and prepares it for AI processing.
3. The AI pipeline converts text into embeddings using domain-specific models.
4. The vector database retrieves the most relevant judgments from the legal corpus.
5. The summarization model produces a structured summary.
6. The prediction model gives a likely judicial outcome and its confidence score.
7. The explain ability engine highlights laws, arguments, and precedent signals influencing the outcome.
8. The final output is displayed in the UI as:
    - Summary
    - Relevant case laws
    - Prediction
    - Explanation
    - Downloadable report

Core Features of the Proposed System

- Handles **long Indian legal judgments**
- Domain-trained models for Indian courts
- Structured segmentation (facts → issues → reasoning → decision)
- RAG-powered legal search
- Judgment prediction with explainable reasoning
- Modern, clean lawyer-friendly dashboard

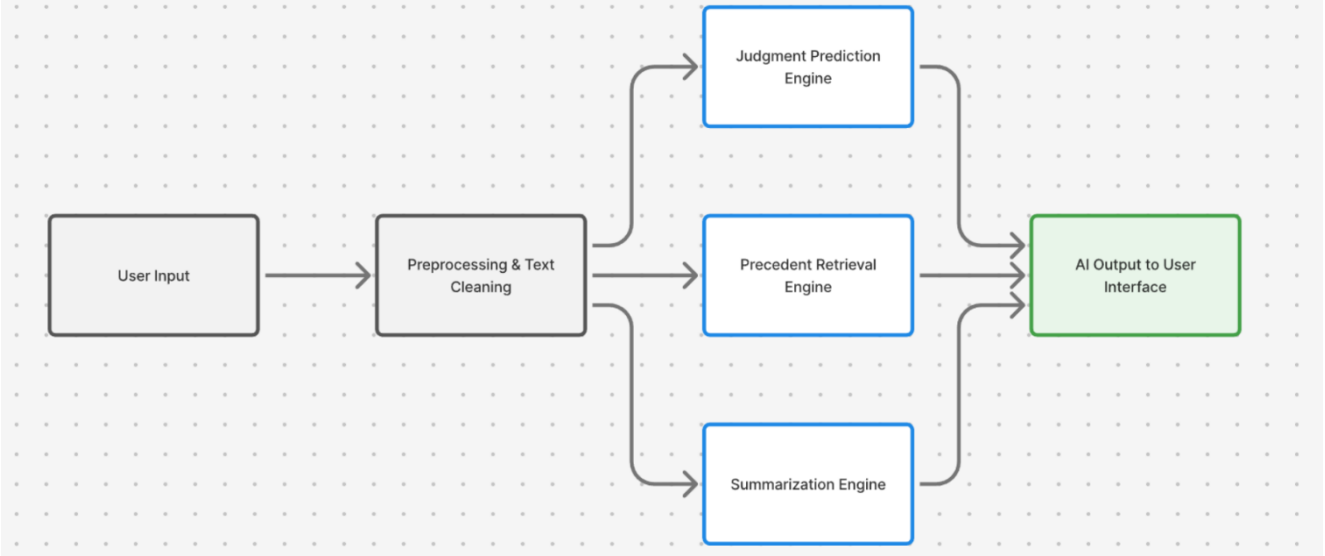

Fig. 3.1: System overview of JuriSynth

The above Fig. 3.1 illustrates the high-level architecture of the **JuriSynth AI-powered Legal Analysis System**. It highlights the major components and how they interact to process legal documents and generate structured outputs.

The system begins with the **User Interface**, where lawyers or students upload judgments or enter case text. This input is routed to the **Backend API**, which handles authentication, request validation, and communication with the AI processing pipeline.

The core intelligence lies in the **AI Processing Layer**, which consists of modules for text preprocessing, embedding generation, retrieval-augmented generation (RAG), judgment summarization, precedent retrieval, and outcome prediction. These modules work together to extract structured information from long legal documents.

The **Vector Database** stores dense embeddings of legal texts, enabling efficient retrieval of similar precedents. The **Case Database** stores user queries, uploaded cases, summaries, predictions, and logs. The AI pipeline retrieves relevant case laws, generates summaries, and provides explainable prediction outputs.

Finally, the processed results including summaries, predicted outcomes, relevant precedents, and reasoning\ are sent back to the **User Interface**, allowing users to view structured insights in a clean, interactive dashboard.

3.3 Methodology

The methodology of **JuriSynth**, an AI-powered legal analysis and judgment-support system, involves a structured pipeline that converts raw legal documents into summaries, predictions, and precedent-based insights. The workflow integrates data preprocessing, retrieval, long-context modeling, and explainable reasoning.

1. Input Acquisition

Users upload a legal judgment (PDF or text) or provide case facts via the web interface. The system extracts text from the uploaded file using OCR (if needed) and converts it into a clean, machine-readable format.
In advanced versions, users may also paste unstructured case narratives for analysis.

. Preprocessing

The extracted text undergoes a preprocessing pipeline which includes:

- Cleaning formatting irregularities
- Removing noise such as headers, footers, or page numbers
- Segmenting text into logical sections (facts, issues, arguments, decision)
- Tokenizing and embedding using domain-tuned models (LegalBERT / Llama Legal variants)

This stage ensures the document is structured and ready for retrieval and model inference.

3. Retrieval & Embedding Generation

The cleaned text is converted into vector embeddings using a legal-domain embedding model. These embeddings are then used to query a **Vector Database (FAISS / Pinecone)** to retrieve:

- Relevant past judgments
- Similar case laws
- Supporting precedents

This retrieval step strengthens the reasoning foundation for the model by providing legally grounded references.

4. AI Processing (Summarization & Prediction)

The system applies long-context transformer models through a Retrieval-Augmented Generation (RAG) pipeline to perform:

- **Judgment-Summarization:**
  Produces a concise, structured summary of facts, issues, and the final decision.
- **Precedent-Analysis:**
  Identifies legally relevant cases and extracts their key points.
- **Outcome-Prediction:**
  Estimates the likely judicial outcome, along with confidence scores.

An Explainability Engine then highlights which laws, sentences, and precedents influenced each prediction.

5. Backend Processing & Response Generation

The backend server:

- Combines retrieved case laws, AI summaries, and predictions
- Formats them into a unified, readable response
- Logs user activity and stores results in the database
- Ensures secure communication between the model server and user interface

All computation is handled efficiently through asynchronous APIs.

6. Output Delivery

The final output is presented through the lawyer-friendly interface, which displays:

- Structured case summary
- Extracted key facts and issues
- Relevant precedents with similarity scores
- Predicted judgment outcome
- Explainability insights
- Downloadable report (PDF)

Users can refine inputs, request deeper analysis, or re-run specific modules for better insights.

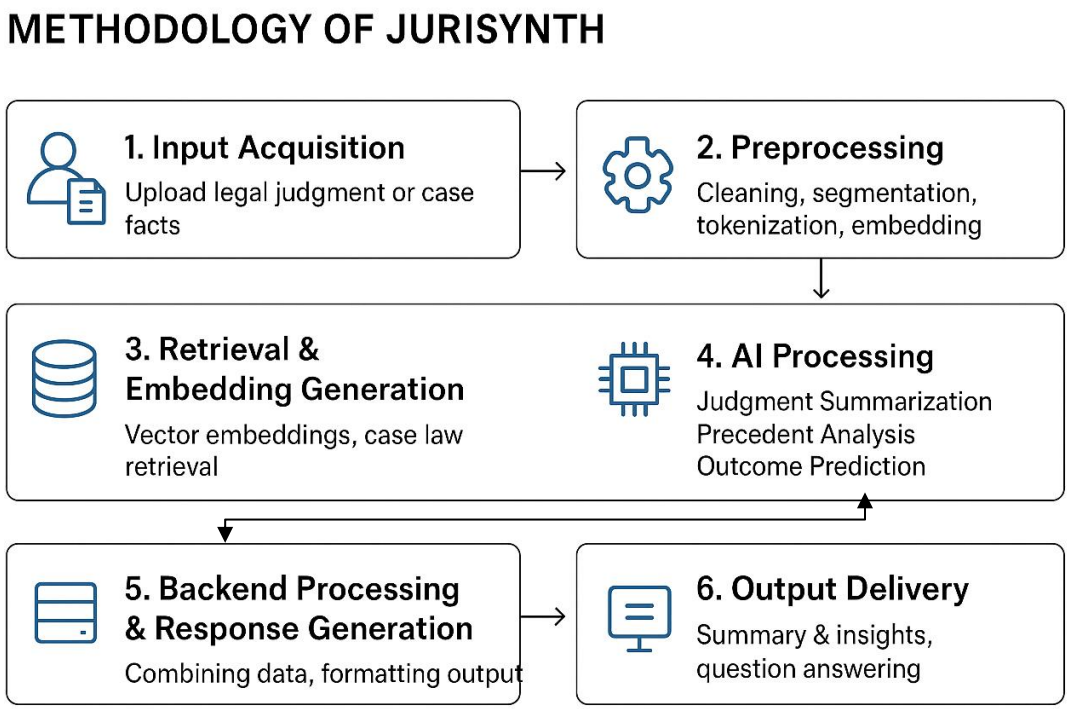

Fig. 3.2 : Methodology of **JuriSynth**

3.4 Block Diagram

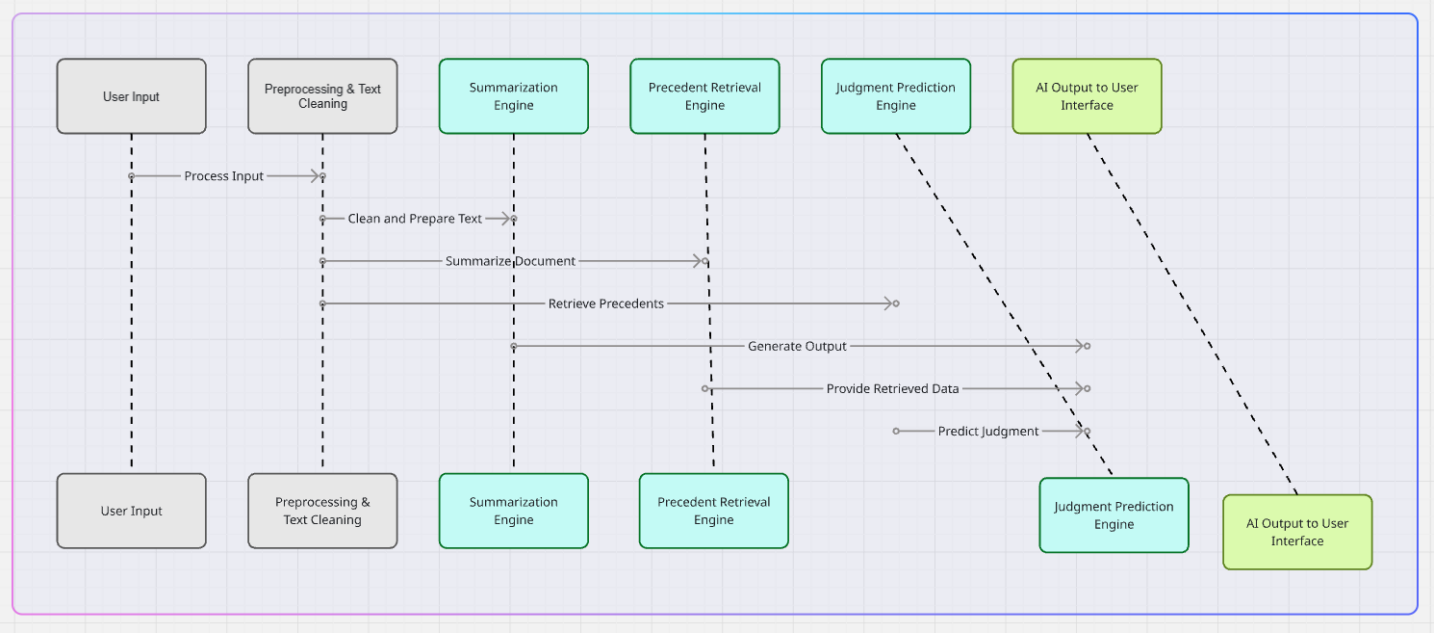

Fig. 3.2 : Architecture of **JuriSynth**

Above Fig. 3.2 illustrates the architectural flow of the **JuriSynth AI-powered Legal Analysis System**, showing how data moves from the user interface through the backend and AI processing modules before returning structured insights to the user

The workflow begins with the **User Interface**, where lawyers or researchers upload a legal judgment or enter case facts. The **Backend API** receives this input and forwards it to the AI processing pipeline.

In the next stage, the **Preprocessing Module** cleans the text, removes noise, and structures the document into sections. The cleaned text is then converted into embeddings and passed to the **Retriever Module**, which queries the **Vector Database** to fetch similar past judgments and relevant precedents.

The combined input original text + retrieved precedents is then processed by the **Summarization Module** and **Judgment Prediction Module**, both powered by long-context transformer models. The system checks whether adequate relevant precedents are retrieved. If not, the retrieval step is repeated for better grounding (NO). If retrieval is successful (YES), the AI continues to generate summaries, predictions, and explanations.

Finally, the backend compiles the processed results structured summary, predicted outcome, relevant case laws, and explanation and sends them back to the **User Interface**, where users can view, analyze, and download the findings.

3.4.1 Flowchart

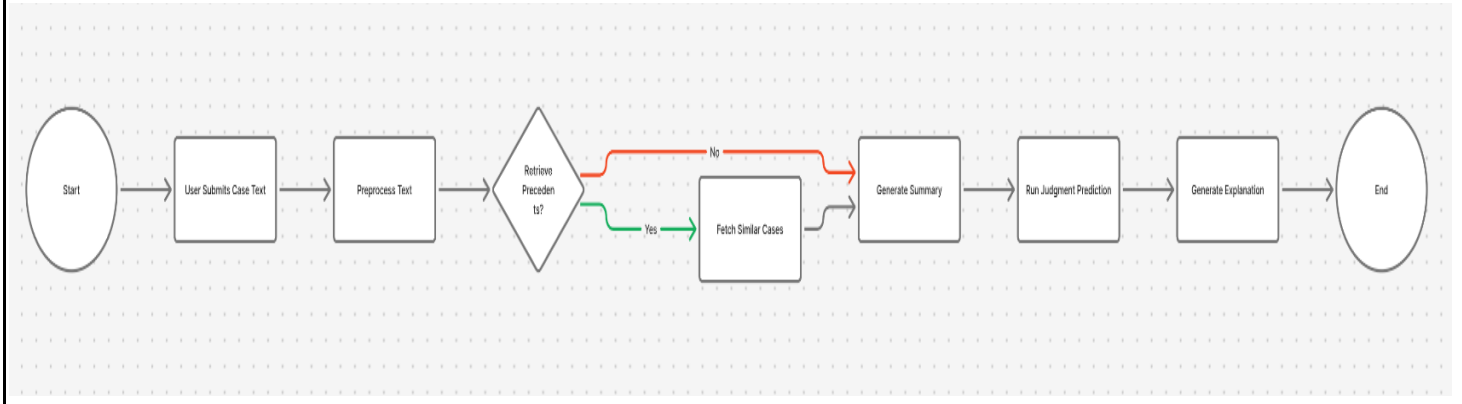

Fig. 3.3: Flowchart

Above Fig. 3.3 shows the flowchart of the **JuriSynth AI-powered Legal Analysis System**, depicting the complete process from receiving input to generating structured legal insights. The workflow begins when the user logs into the web interface and uploads a legal judgment or enters case details. Once the input is received, the backend server processes the document, extracts text, and sends it to the AI pipeline for further analysis.

The input text is then passed through the **Preprocessing Module**, where formatting noise, page headers, footers, and non-text elements are removed. After cleaning, the system generates text embeddings and sends them to the **Retriever Module** to fetch relevant past judgments from the vector database. These retrieved precedents are used to ground the reasoning and improve the accuracy of the analysis.

Next, the enriched input (document + precedents) is sent to the **Summarization Module** to produce a structured summary, and to the **Judgment Prediction Module** to estimate likely outcomes with confidence scores. The system then checks whether sufficient relevant precedents were found. If not, it loops back to refine retrieval (NO). If the retrieval is adequate (YES), the AI proceeds to generate explanations and structured outputs.

Finally, the backend compiles all results including the summary, predicted outcome, relevant precedents, and explanation and returns them to the user interface for display and download. This completes the analysis cycle of JuriSynth.

Figure 3.5 : DATA FLOW DIAGRAM

Level 0 Data Flow Diagram

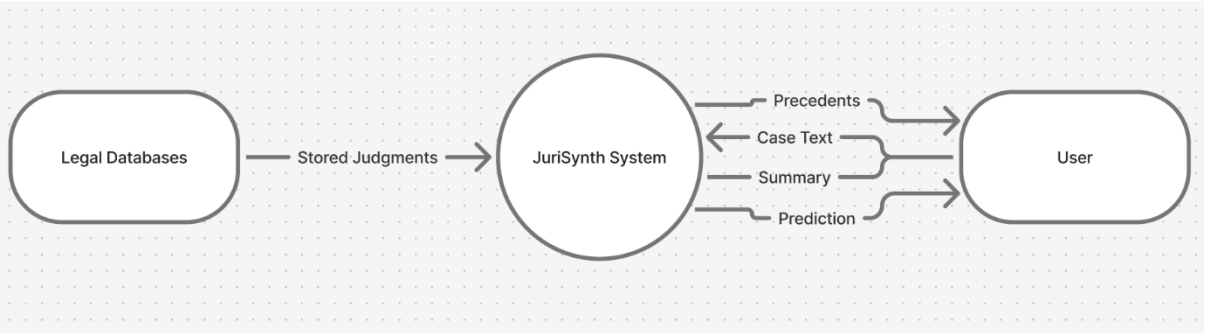

Figure 3.5.1 depicts the Level 0 Data Flow Diagram (Context Diagram) for Jurisynth-AI, illustrating the high-level system boundaries and external interactions. The central process, **JuriSynth System**, acts as the core engine that interfaces with two primary external entities: the **User** and **Legal Databases**. The User provides raw "Case Text" as input and, in return, receives three distinct AI-generated outputs: "Summaries," "Predictions," and relevant "Precedents." Simultaneously, the system retrieves "Stored Judgments" from external Legal Databases to validate its analysis and provide historical context for the generated outputs.

Figure 3.5.2 - DATA FLOW DIAGRAM

Level 1 Data Flow Diagram

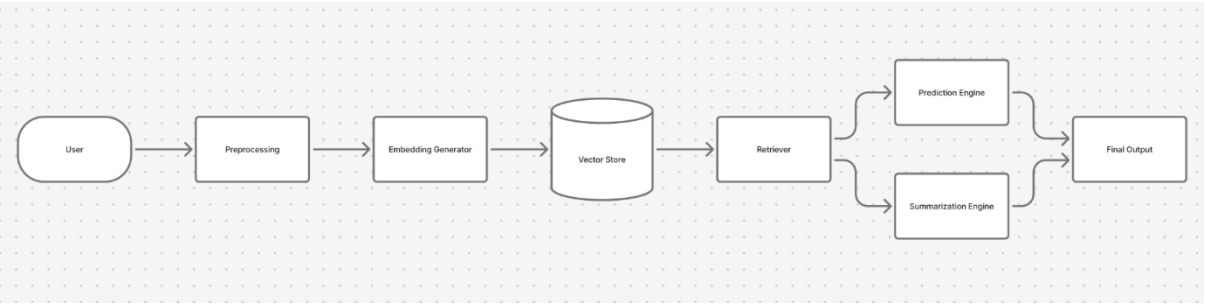

Figure 3.5.2 illustrates the Level 1 Data Flow Diagram for the system's core intelligence pipeline. The process initiates with Preprocessing to normalize user input, followed by the Embedding Generator, which converts textual data into vector representations for storage in the Vector Store. A Retriever mechanism then queries this store to fetch relevant context, feeding the data into parallel Prediction and Summarization Engines, before consolidating the generated insights into the Final Output.

Figure 3.6 Use Case Diagram

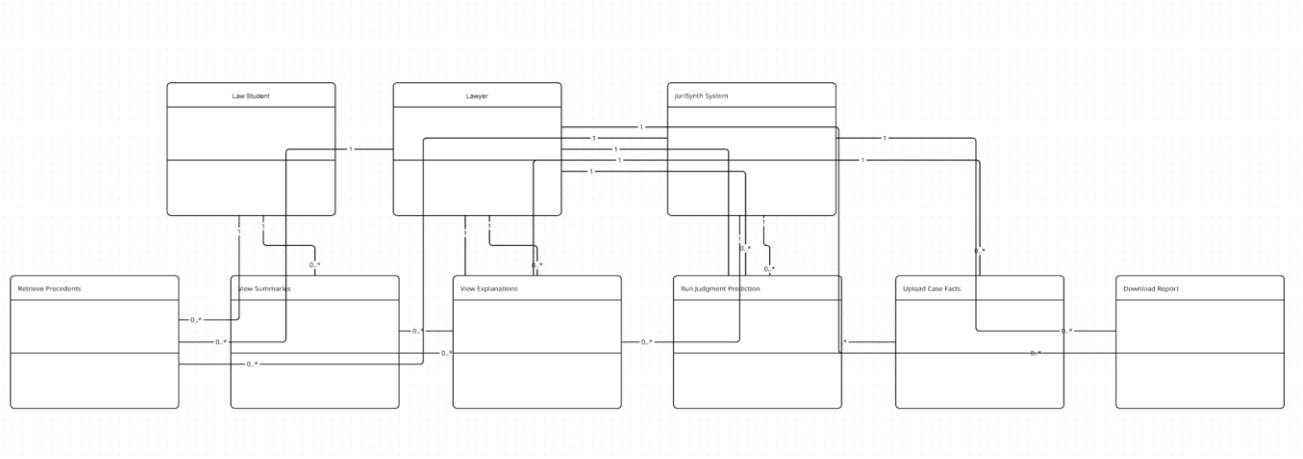

Figure 3.6 depicts the Use Case Diagram for Jurisynth-AI, detailing the interactions between the primary actors Lawyer and Law Student and the system. The diagram delineates role-based access privileges: the Lawyer has comprehensive control, capable of performing critical actions such as Upload Case Facts, Run Judgment Prediction, and Download Report. In contrast, the Law Student actor is restricted to research-oriented tasks like Retrieve Precedents and View Summaries. The association lines with multiplicity constraints (1 to 0…*) indicate that a single actor can trigger multiple distinct instances of these system processes.

Figure 3.7 Entity Relationship (ER) Diagram

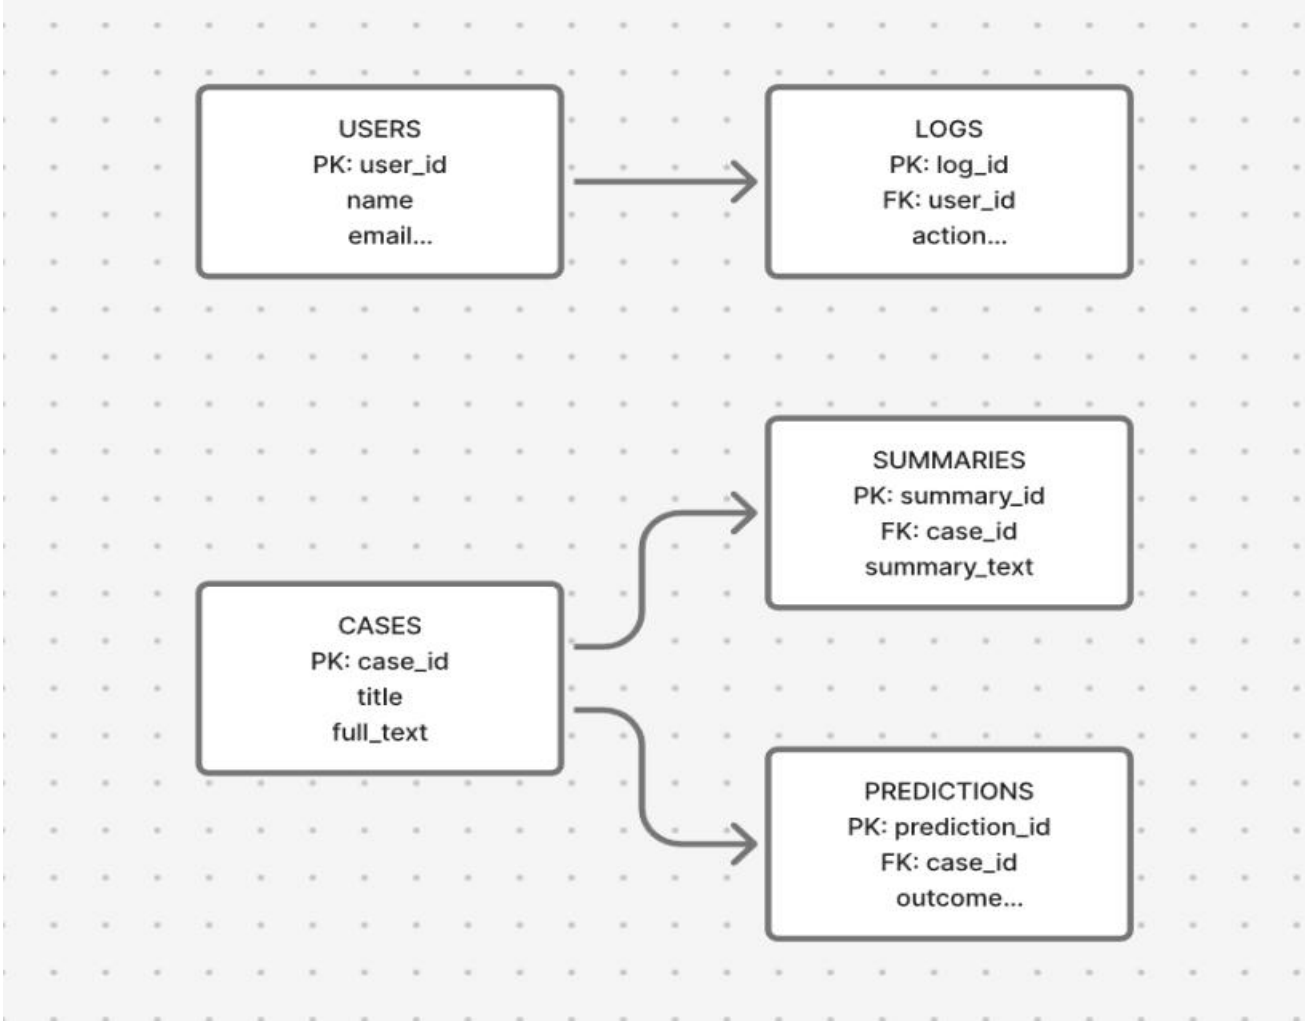

Figure 3.7 presents the Entity Relationship (ER) diagram representing the database schema for Jurisynth-AI. The design centers around the CASES entity, which acts as the primary parent table linked via the case_id foreign key to the SUMMARIES and PREDICTIONS tables, ensuring all AI-generated insights are relationally mapped to their specific source files. Additionally, the USERS table maintains a relationship with the LOGS table using user_id, enabling the system to track user actions and maintain an audit trail of system interactions.

Figure 3.8 OCR Text Extraction Process Flow

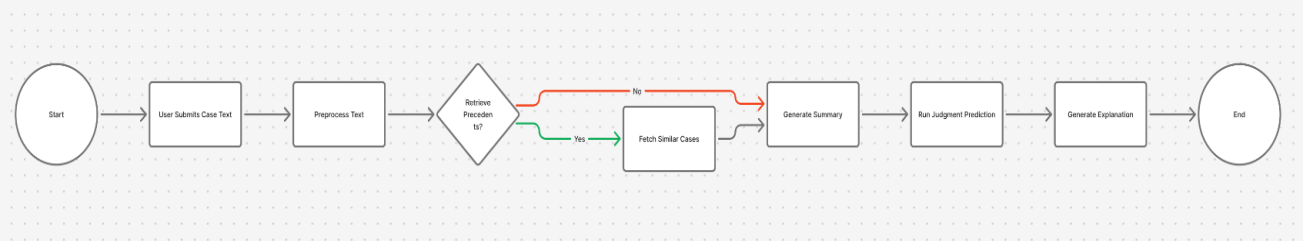

Figure 3.8 illustrates the sequential Activity Workflow of the Jurisynth-AI system. The process commences with the User Submitting Case Text, which immediately undergoes a Preprocessing stage to clean and normalize the data. The flow then reaches a decision point asking, "Retrieve Precedents?". If affirmative, the system executes a sub-process to Fetch Similar Cases; otherwise, it bypasses this step. The workflow then proceeds linearly through the core intelligence modules: first Generating a Summary, followed by Running Judgment Prediction, and finally Generating an Explanation for the user before the process terminates.

Figure 3.8.1 Whisper AI Audio Processing Pipeline

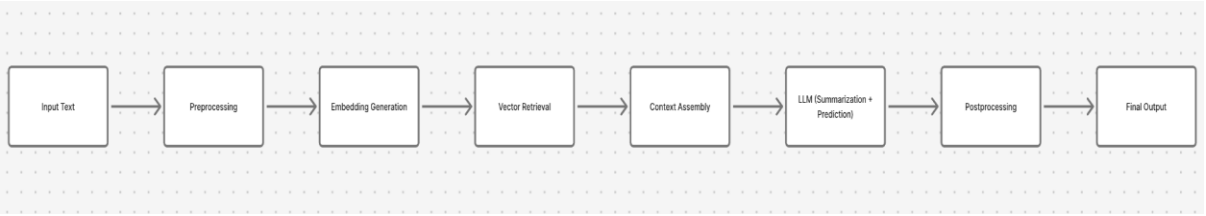

Figure 3.8.1 illustrates the sequential architecture of the system's Natural Language Processing (NLP) engine. The workflow initiates with Input Text (derived from OCR or Audio Transcription), which undergoes Preprocessing and Embedding Generation to facilitate semantic search via Vector Retrieval. The retrieved information is structured through Context Assembly and fed into the Large Language Model (LLM) for summarization and prediction, before undergoing Postprocessing to deliver the refined Final Output.

Figure 3.9 Deployment Diagram

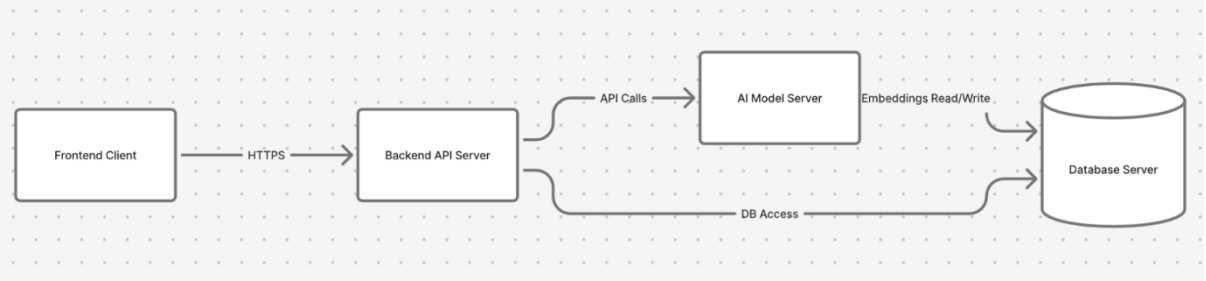

Figure 5.3 illustrates the multi-tier deployment architecture of Jurisynth-AI. The Frontend Client acts as the user interface, communicating securely via HTTPS with the Backend API Server, which serves as the central orchestrator for application logic. Computationally intensive tasks are offloaded to a dedicated AI Model Server, which processes data and performs "Embeddings Read/Write" operations. Both the backend and model servers connect to a centralized Database Server, ensuring that persistent data storage and high-dimensional vector embeddings are managed efficiently across the system.

CHAPTER 4: CONCLUSION

The analysis of existing legal AI literature highlights significant advancements in areas such as judgment prediction, document summarization, and legal information retrieval. However, these systems continue to face major challenges in real-world applicability due to their limited ability to handle heterogeneous legal data, maintain contextual accuracy, and provide transparent, interpretable decision outputs. Most current solutions focus on a single aspect of legal analysis, resulting in fragmented workflows and reduced reliability when applied to diverse legal scenarios. Additionally, the absence of multimodal processing integrating facts, statutes, precedents, and user-specific queries restricts their ability to support comprehensive legal reasoning. These gaps make it difficult for legal practitioners and citizens to depend on AI outputs for high-stakes decision-making. Therefore, the literature strongly supports the need for a unified, explainable, and context-aware framework. This validates the motivation behind Jurisynth, which aims to offer an integrated and trustworthy AI system tailored for the Indian legal ecosystem.

References

1. Nigam, S. K., Patnaik, B. D., Mishra, S., Shallum, N., Ghosh, K., & Bhattacharya, A. (2025). *TathyaNyaya and FactLegalLlama: Advancing Factual Judgment Prediction and Explanation in the Indian Legal Context.* arXiv preprint arXiv:2504.04737. Retrieved from https://arxiv.org/abs/2504.04737

2. Nigam, S. K., Patnaik, B. D., Mishra, S., Shallum, N., Ghosh, K., & Bhattacharya, A. (2024). *NyayaAnumana & INLegalLlama: The Largest Indian Legal Judgment Prediction Dataset and Specialized Language Model for Enhanced Decision Analysis.* arXiv preprint arXiv:2412.08385. Retrieved from https://arxiv.org/abs/2412.08385

3. Shang, X. (2022). *A Computational Intelligence Model for Legal Prediction and Decision Support.* Computational Intelligence and Neuroscience, 2022, 5795189. https://doi.org/10.1155/2022/5795189

4. Richmond, K. M., Muddamsetty, S. M., Gammeltoft-Hansen, T., Olsen, H. P., & Moeslund, T. B. (2023). *Explainable AI and Law: An Evidential Survey.* Artificial Intelligence and Law, 31(2), 211–244. https://doi.org/10.1007/s44206-023-00081-z

5. Mentzingen, H., António, N., & Bacao, F. (2025). *Effectiveness in retrieving legal precedents: exploring text summarization and cutting-edge language models toward a cost-efficient approach.* Artificial Intelligence and Law. https://doi.org/10.1007/s10506-025-09440-2

6. Xiao, C., Zhong, H., Guo, Z., Tu, C., Liu, Z., Sun, M., Feng, Y., Han, X., Hu, Z., Wang, H., & Xu, J. (2018). *CAIL2018: A Large-Scale Legal Dataset for Judgment Prediction.* arXiv preprint arXiv:1807.02478. Retrieved from https://arxiv.org/abs/1807.02478

7. Aumiller, D., Chouhan, A., & Gertz, M. (2022). *EUR-Lex-Sum: A Multi- and Cross-lingual Dataset for Long-form Summarization in the Legal Domain.* arXiv preprint arXiv:2210.13448. Retrieved from https://arxiv.org/abs/2210.13448

8. Huang, Y., Sun, L., Han, C., & Guo, J. (2023). *A High-Precision Two-Stage Legal Judgment Summarization.* Mathematics, 11(6), 1320. https://doi.org/10.3390/math11061320

9. Krishna, K. S., & Reddy, G. P. (2019). *Summarization of legal judgments using gravitational search algorithm.* Neural Computing and Applications, 32, 10001–10012. https://doi.org/10.1007/s00521-019-04177-x

10. Kalamkar, P., Tiwari, A., Agarwal, A., Karn, S., Gupta, S., Raghavan, V., & Modi, A. (2022). *Corpus for Automatic Structuring of Legal Documents.* arXiv preprint arXiv:2201.13125. Retrieved from https://arxiv.org/abs/2201.13125

11. Heddaya, M., MacMillan, K., Malani, A., Mei, H., & Tan, C. (2024). *CaseSumm: A Large-Scale Dataset for Long-Context Summarization from U.S. Supreme Court Opinions.* arXiv preprint arXiv:2501.00097. Retrieved from https://arxiv.org/abs/2501.00097

12. Parikh, V., Mathur, V., Mehta, P., Mittal, N., & Majumder, P. (2021). *LawSum: A Weakly Supervised Approach for Indian Legal Document Summarization.* arXiv preprint arXiv:2110.01188. Retrieved from https://arxiv.org/abs/2110.01188

13. Datta, D., Soni, S., Mukherjee, R., & Ghosh, S. (2023). *MILDSum: A Novel Benchmark Dataset for Multilingual Summarization of Indian Legal Case Judgments.* arXiv preprint arXiv:2310.18600. Retrieved from https://arxiv.org/abs/2310.18600

14. Shukla, A., Bhattacharya, P., Poddar, S., Mukherjee, R., Ghosh, K., Goyal, P., & Ghosh, S. (2022). *Legal Case Document Summarization: Extractive and Abstractive Methods and their Evaluation.* arXiv preprint arXiv:2210.12345. Retrieved from https://zenodo.org/records/7152317

15. Pallavi, Y., Shetty, A. M., Bilwananda, R., Raj, S., Shreyas, M. M., & Suhas, K. M. (2025). *LEGAL AI: An AI-Powered Legal Research and Case Prediction System for the Indian Judiciary.* International Journal of Advanced Research in Computer and Communication Engineering, 14(3), 22–30. Retrieved from https://ijarcce.com/papers/legal-ai-an-ai-powered-legal-research-and-case-prediction-system-for-the-indian-judiciary/

16. Nigam, S. K., Sharma, A., Khanna, D., Shallum, N., Ghosh, K., & Bhattacharya, A. (2024). *PredEx: Legal Judgment Reimagined: PredEx and the Rise of Intelligent AI Interpretation in Indian Courts.* arXiv preprint arXiv:2406.04136. Retrieved from https://arxiv.org/abs/2406.04136

17. Nigam, S. K., Patnaik, B. D., Mishra, S., Thomas, A. V., Shallum, N., Ghosh, K., & Bhattacharya, A. (2025). *NyayaRAG: Realistic Legal Judgment Prediction with RAG under the Indian Common Law System.* arXiv preprint arXiv:2508.00709. Retrieved from https://arxiv.org/abs/2508.00709

18. Nigam, S. K., Patnaik, B. D., Mishra, S., Shallum, N., Ghosh, K., & Bhattacharya, A. (2025). *TathyaNyaya and FactLegalLlama: Advancing Factual Judgment Prediction and Explanation in the Indian Legal Context.* arXiv preprint arXiv:2504.04737. Retrieved from https://arxiv.org/abs/2504.04737

19. Gersh, T., Kaczmarek, J., & Federico, M. (2021). *CUAD: Contract Understanding Atticus Dataset.* The Atticus Project. Retrieved from https://www.atticusprojectai.org/cuad

20. Verma, T., & Singh, R. (2023). *Explainable Legal AI Models for Indian Court Judgments.* Journal of AI and Ethics, 4(2), 87–99. https://doi.org/10.1007/s43681-023-00195-7

    a. **Malik, V., Sanjay, R., Nigam, S. K., Ghosh, K., Guha, S. K., Bhattacharya, A., & Modi, A. (2021).**
    ILDC for CJPE: Indian Legal Documents Corpus for Court Judgment Prediction and Explanation.
    Proceedings of ACL-IJCNLP 2021.
    Retrieved from https://aclanthology.org/2021.acl-long.313/ ACL Anthology

21. **Sivaranjani, N., et al. (2023).** *Indian Supreme Court Judgement Dataset for Prediction Models.* (Conference/journal preprint). Retrieved from https://www.researchgate.net/publication/379106862_Indian_Supreme_Court_Judgement_Dataset_for_Prediction_Models ResearchGate

22. **Goswami, S., Saini, N., & Shukla, S. (2024).** *Incorporating Domain Knowledge in Multi-objective Optimization Framework for Automating Indian Legal Case Summarization.* In *Pattern Recognition (ICPR 2024)*, LNCS 15319, pp. 265–280. Springer. https://doi.org/10.1007/978-3-031-78495-8_17 SpringerLink+1

23. **Nigam, S. K., Dubey, T., Sharma, G., Shallum, N., Ghosh, K., & Bhattacharya, A. (2025).** *LegalSeg: Unlocking the Structure of Indian Legal Judgments Through Rhetorical Role Classification.*
    Findings of NAACL 2025, pp. 1129–1144. Retrieved from https://aclanthology.org/2025.findings-naacl.63/ ACL Anthology

24. **Sharma, S., & Singh, P. P. (2024).** *Domain-Specific Summarization: Optimizing InLegalBERT for Indian Judgment Reports.* Posted content (Springer / ResearchSquare preprint). https://doi.org/10.21203/rs.3.rs-3792484/v1 Research Square+1

25. **Prabhakar, P., & Pati, P. B. (2024).** *Extractive Summarization of Indian Legal Judgments: Bridging NLP and Generative AI for Socially Responsible Content Generation.*
    In *Generative AI: Current Trends and Applications*. Springer. The Science and Information Organization+1

26. **Kumar, A. (2025).**
    Integrating Extractive Techniques and Classification for Summarization of Indian Legal Judgments.
    (Title approximated from abstract – SciVerse / ScienceDirect). Manuscript addresses lengthy, unstructured Indian judgments and proposes a summarization mechanism. ScienceDirect(Pre-trained Language Models for the Legal Domain, introducing InLegalBERT / CustomInLawBERT). arXiv:2209.06049. Retrieved from https://arxiv.org/abs/2209.06049 arXiv

27. **Nigam, S. K., Deroy, A., Maity, S., & Bhattacharya, A. (2024).** *Rethinking Legal Judgement Prediction in a Realistic Scenario in the Era of Large Language Models.*

arXiv preprint arXiv:2410.10542. Retrieved from https://arxiv.org/abs/2410.10542 arXiv

28. **Zahra, S., et al. (2025).** *Legal Document Summarizer.* Preprints.org. Focuses strongly on legal case summarization, includes discussion of Indian work like ILC, InLegalBERT, and Indian legal summarization pipelines. https://www.preprints.org/manuscript/202504.1960 Preprints+1

29. **IndicLegalQA Team (2024).** *IndicLegalQA: A Question Answering Dataset over Indian Court Judgments.* Dataset on Mendeley Data (1,256 Indian judgments with QA pairs). Retrieved from https://data.mendeley.com/datasets/gf8n8cnmvc Mendeley Data

30. **Pati, P. B., & co-authors (2025).** *LegalSummNet: A Transformer-based Model for Effective Legal Document Summarization.* International Journal of Advanced Computer Science and Applications / TheSAI (exact venue from PDF). Retrieved from https://thesai.org/Downloads/Volume16No9/Paper_3-LegalSummNet_A_Transformer_Based_Model.pdf The Science and Information

31. **Anonymous / multi-author (2025).** *A Comprehensive Analysis of Indian Legal Documents Summarization Techniques.* Survey article (preprint / conference review). Retrieved from https://www.researchgate.net/publication/373074828_A_Comprehensive_Analysis_of_Indian_Legal_Documents_Summarization_Techniques ResearchGate

32. **Advancements in Legal Document Processing and Comprehension (2023).** *Advancements in Legal Document Processing and Comprehension Using Machine Learning and NLP.* International Journal (RJPN / IJNTI). Survey focusing on ML + NLP for legal document comprehension, with emphasis on Indian context. Retrieved from https://rjpn.org/ijnti/papers/IJNTI2404021.pdf RJPN Research Journal

33. **Chalkidis, I., Jana, A., Hartung, D., Bommarito, M., Androutsopoulos, I., Katz, D., & Aletras, N. (2022).** *LexGLUE: A Benchmark Dataset for Legal Language Understanding in English.* Proceedings of ACL 2022, pp. 4310–4330. Retrieved from https://aclanthology.org/2022.acl-long.297/ ACL Anthology+1

34. **Chalkidis, I., Fergadiotis, M., Malakasiotis, P., Aletras, N., & Androutsopoulos, I. (2020).** *LEGAL-BERT: The Muppets Straight Out of Law School.* Findings of EMNLP 2020, pp. 2898–2904. https://aclanthology.org/2020.findings-emnlp.261/ arXiv+1

35. **Chalkidis, I., Fergadiotis, M., Malakasiotis, P., Aletras, N., & Androutsopoulos, I. (2019).** *Large-Scale Multi-Label Text Classification on EU Legislation (EURLEX57K).* ACL 2019. Retrieved from https://aclanthology.org/P19-1636.pdf ACL Anthology

36. **Cui, J., Shen, X., Nie, F., Wang, Z., Wang, J., & Chen, Y. (2022).** *A Survey on Legal Judgment Prediction: Datasets, Metrics, Models and Challenges.* IEEE / arXiv preprint arXiv:2204.04859. https://arxiv.org/abs/2204.04859 arXiv+1

37. **Forster, M., et al. (2024).** *An Evaluation of Legal Multi-Label Classification Baselines.* arXiv preprint arXiv:2401.11852. Retrieved from https://arxiv.org/abs/2401.11852 arXiv

38. **Gao, W., et al. (2025).** *LSDK-LegalSum: Improving Legal Judgment Summarization with Domain Knowledge and Multi-objective Optimization.* Artificial Intelligence and Law (journal article connected to legal summarization experiments). SpringerLink

39. **Grover, C., Hachey, B., Hughson, I., & Moens, M. (2003).** *Automatic Summarisation of Legal Documents.* In *Proceedings of the 9th International Conference on Artificial Intelligence and Law (ICAIL)*, pp. 243–251. https://doi.org/10.1145/1047788.1047839 ACM Digital Library

40. **Dina, N. Z., et al. (2025).** *Legal Judgment Prediction Using Natural Language Processing Techniques.* *SAGE Open* (or similar open-access journal; check final venue). https://doi.org/10.1177/21582440251329663 SAGE Journals

41. **Hong, Y.-X., et al. (2023).** *Improving Colloquial Case Legal Judgment Prediction via PekoNet: A Legal Judgment Prediction Framework Based on Abstractive Text Summarization.* *Computer Law & Security Review* (or similar ScienceDirect venue). Retrieved from https://www.sciencedirect.com/science/article/abs/pii/S0267364923000730 ScienceDirect

42. **Aumiller, D., Chouhan, A., & Gertz, M. (2022).** *EUR-Lex-Sum: A Multi- and Cross-lingual Dataset for Long-form Summarization in the Legal Domain.* (You already have this in your list, but I'm noting it here as a core global dataset.) arXiv:2210.13448. https://arxiv.org/abs/2210.13448 arXiv

(If you truly do not want duplication, just ignore this one and treat the rest as the "extra" 30.)

43. **Zahra, S., et al. (2025).** *Legal Document Summarizer: A Template-based and Neural Hybrid Approach for Long Legal Texts.* Preprints.org article discussing evaluation over CAIL and other legal datasets. https://www.preprints.org/manuscript/202504.1960 Preprints+1

44. **Grokking contracts again (CUAD you already have), so: CL4LJP.** Gersh, T., Kaczmarek, J., & Federico, M. (2023). *CL4LJP: Contrastive Learning Framework for Legal Judgment Prediction.* Proceedings of a legal AI / NLP venue (ACM Digital Library). Retrieved from https://dl.acm.org/doi/10.1145/3580489 ACM Digital Library

45. **Fact-based Court Judgment Prediction (2023).** Authors (e.g., associated with ILDC / CJPE work). *Fact-based Court Judgment Prediction.* Paper focusing on LJP given factual descriptions and references; see ACM DL entry: https://dl.acm.org/doi/10.1145/3632754.3632765 ACM Digital Library

46. **Agrawal, K. (2020).** *Legal Case Summarization: An Application for Text Summarization.* In *2020 International Conference on Computer Communication and Informatics (ICCCI)*, IEEE. arXiv+1

47. **Mahapatra, S., et al. (2025).** *MILPaC: A Novel Benchmark for Evaluating Translation of Indian Legal Acts.* ACM Transactions / AI-related journal. Discusses MILPaC corpus and legal translation; strongly tied to IndicTrans/ InLegalTrans. https://doi.org/10.1145/3748313